

Stochastic Modeling

Philippe Chevalier
UCLouvain

Section 2: Finite State Markov Chains



Introduction

definition

A Markov Chain is a stochastic process with fixed intervals $\{X_n, n \geq 0\}$ such that each random variable $X_n, n \geq 1$ depends on the past only through the most recent random variable X_{n-1}

$$\begin{aligned} P[X_n = j | X_{n-1} = i, X_{n-2} = k, \dots, X_0 = m] &= P[X_n = j | X_{n-1} = \\ &= P_{ij} \end{aligned}$$

The random variable X_n is called the state of the Markov Chain


The set of possible sample values for the states lie in a countable set (of finite size in this section)

We will generally call the possible states $\{1, 2, \dots, J\}$.

When the transition probabilities depend on the period ($p_{ij}(n)$), we have a non-homogeneous Markov Chain (not studied in this course)

Two representations

Graphical



Matrix

$$P = [P_{ij}] = \begin{bmatrix} P_{11} & P_{12} & \cdots & P_{1J} \\ P_{21} & P_{22} & \cdots & P_{2J} \\ \vdots & & \ddots & \vdots \\ P_{J1} & P_{J2} & \cdots & P_{JJ} \end{bmatrix}$$

Definition

A state j is accessible from i (abbreviated as $i \rightarrow j$) if there is a walk in the graph from i to j

- ▶ We denote $P_{ij}^n = P[X_n = j | X_0 = i]$
- ▶ $i \rightarrow j$ if and only if $P_{ij}^n > 0$ for some n .

Definition

Two distinct states i and j communicate (abbreviated as $i \leftrightarrow j$) if i is accessible from j and j is accessible from i
 $i \leftrightarrow i$ by definition

Note that if $i \leftrightarrow j$ and $m \leftrightarrow j$ then $i \leftrightarrow m$

Classification of states

Definition

A class C of states is a non-empty set of states such that for each $i \in C$, each state $j \neq i$ satisfies $j \in C$ if $i \leftrightarrow j$ and $j \notin C$ if $j \not\leftrightarrow i$.

- ▶ A state i is *recurrent* if it is accessible from all states that are accessible from i . (i is recurrent if $i \rightarrow j$ implies that $j \rightarrow i$).
- ▶ A transient state is a state that is not recurrent.
- ▶ All states from a same class will be of the same type.
- ▶ A finite state Markov Chain has at least one recurrent class.

Definition

The period of a state i , denoted $d(i)$, is the greatest common divisor (gcd) of those values of n for which $P_{ii}^n > 0$.

If the period is 1, the state is *aperiodic*, and if the period is 2 or more, the state is *periodic*.

- ▶ All the states of a class will have the same periodicity
- ▶ If a class has a period $d > 1$, then there exists a partition C_1, C_2, \dots, C_d of the states of this class in d subclasses such that all the transitions from a state of C_n go to a state of class C_{n+1} and all transitions from C_d go to a state of C_1
- ▶ If $j \in C_n$ and $p_{jk} > 0$ then $k \in C_{n+1}$
- ▶ A class that is aperiodic and recurrent is called *ergodic*

Computing transition probabilities

$$P[X_{n+2} = j | X_n = i] = P_{ij}^2 = \sum_{k=1}^J P_{ik} P_{kj}$$
$$\Rightarrow P^2 = P.P \quad \text{et} \quad P^n = \underbrace{P.P.P \dots P}_n$$

More generally:

$$P_{ij}^{m+n} = \sum_{k=1}^J P_{ik}^m P_{kj}^n$$

$$P^{m+n} = P^m.P^n$$

Convergence

- ▶ A matrix P will be called a stochastic matrix iff the matrix is square, non-negative and each row sums to 1.
- ▶ What would be the limit of P^n when n tends to infinity?
- ▶ If P is ergodic, we would expect all rows to converge towards the same value, let us call this row vector π , then
$$P^\infty = P^\infty P \Rightarrow \pi = \pi P$$
We must also impose that the values of the vector π sum to 1.
- ▶ When does this system have a solution?
When is it unique?
When does P^∞ converge?

A bit of Matrix theory

- ▶ Note that because P is a stochastic matrix we have : $Pe = e$
- ▶ This implies that 1 is an eigenvalue of the matrix P
- ▶ The solution π is a left eigenvector for the eigenvalue 1
- ▶ The number of linearly independent solutions corresponds to the multiplicity of the eigenvalue 1
 - ▶ There will be one independent solution for each recurrent class of P
- ▶ If P is aperiodic then $\lim_{n \rightarrow \infty} P^n = e\pi$
- ▶ else π will be the average over the different subclasses

Tropical weather

- ▶ A little tired with Belgian weather, you decide to go for a vacation in the tropics ...
- ▶ After some research you find some reliable statistics about the weather there. It appears that the weather follows a clear pattern:
 - ▶ If today is sunny, there is an 80% chance of a sunny day tomorrow, else it will be cloudy.
 - ▶ If today is cloudy, there is a 25% chance that it might rain tomorrow, a 25% chance that it stays cloudy, else the weather will be sunny again.
 - ▶ If it rains, there is a 50% chance that the rain lasts for one day else the rain lasts for 2 days, when the rain finishes, the weather is always sunny.
- ▶ What is the probability of a sunny day on your arrival?
How many days of sunshine would you expect during a week?

Find the Stationary probabilities

$$P = \begin{pmatrix} 0,1 & 0,9 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0,4 & 0 & 0,4 & 0 & 0 & 0,2 & 0 & 0 & 0 & 0 \\ 0,25 & 0 & 0,15 & 0,2 & 0 & 0 & 0,2 & 0 & 0 & 0,2 \\ 0 & 0 & 0 & 0,25 & 0,75 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0,75 & 0,25 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0,5 & 0,5 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0,25 & 0,25 & 0 & 0,5 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0,5 & 0 & 0 & 0,5 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0,25 & 0,75 & 0 & 0 \end{pmatrix}$$

Markov chains with rewards

Let r_i be a reward associated with state i

We could also define r_{ij} as the reward associated to a transition between states i and j , the expected gain for each passage through state i is then:

$$r_i = \sum_j P_{ij} r_{ij}$$


For an ergodic chain, we observe that g the average reward per period will be

$$g = \sum_i r_i \pi_i$$

Transient rewards

We will see that the transient rewards are very interesting to understand the behaviour of Markov chains.

To motivate this, let us use the following example:



Expected reward over multiple transitions

Let X_m be the state at time m and let $R_m = R(X_m)$ be the reward at time m , i.e., if the sample value of X_m is i , then r_i is the sample value of R_m .

Conditional on $X_m = i$, the aggregate expected reward $v_i(n)$ over n periods from X_m to X_{m+n-1} is

$$\begin{aligned}v_i(n) &= E[R(X_m) + R(X_{m+1}) + \cdots + R(X_{m+n-1}) | X_m = i] \\&= r_i + \sum_j P_{ij}r_j + \cdots + \sum_j P_{ij}^{n-1}r_j\end{aligned}$$

In vector notation:

$$\mathbf{v}(n) = \mathbf{r} + [P]\mathbf{r} + \cdots + [P^{n-1}]\mathbf{r} = \sum_{h=0}^{n-1} [P^h]\mathbf{r}$$

Assuming the Markov chain is an *ergodic unichain* (has a single ergodic class with possibly some transient classes),

$$\lim_{n \rightarrow \infty} [P^n] = e\pi$$


$$\Rightarrow \lim_{n \rightarrow \infty} [P^n]r = e\pi r = ge$$

So when n grows, the expected reward per period converges to g .
The transient effect can thus be evaluated by:

$$\lim_{n \rightarrow \infty} (v(n) - nge) = \lim_{n \rightarrow \infty} \sum_{h=0}^{n-1} [P^h - e\pi]r$$

From the properties of P we can show that this limit exist. It is called the relative-gain vector and denoted by w .

Example illustration



n	$v_1(n)$	$v_2(n)$
1	0	1
2	0.01	1.99
4	0.0592	3.9408
10	0.4268	9.5732
40	6.1425	33.8575
100	28.3155	71.6845
400	175.007	224.9923

The relative gain vector w

The relative gain vector w can be computed by solving the following equations

$$w + ge = [P]w + r \quad \text{and} \quad \pi w = 0$$

Proof:



$$\begin{aligned}[P]w &= \lim_{n \rightarrow \infty} \sum_{h=0}^{n-1} [P^{h+1} - e\pi]r = \lim_{n \rightarrow \infty} \sum_{h=1}^n [P^h - e\pi]r \\ &= \lim_{n \rightarrow \infty} \left(\sum_{h=0}^n [P^h - e\pi]r \right) - [P^0 - e\pi]r \\ &= w - [P^0 - e\pi]r \\ &= w - r + ge\end{aligned}$$

Note that this linear system always admits a solution, the normalization constraint on w is arbitrary.

Markov Decision Processes

Suppose that in each state i , we can choose between K^i different possibilities with rewards $r_i^{(1)}, r_i^{(2)}, \dots, r_i^{(K_i)}$ with the corresponding transition probabilities $\{P_{ij}^{(1)}, 1 \leq j \leq J\}, \{P_{ij}^{(2)}, 1 \leq j \leq J\}, \dots, \{P_{ij}^{(K_i)}, 1 \leq j \leq J\}$

Example :



- ▶ Can we find an optimal stationary policy?
- ▶ Can we find an optimal dynamic policy?

Dynamic Programming algorithm for the dynamic optimal policy

Let $v(0)$ be the final reward, 1 period before the end we can compute the optimal action by:

$$v_i^*(1) = \max_k \{ r_i^{(k)} + \sum_j P_{ij}^{(k)} v_j(0) \}$$

The expected global reward with 2 periods to go is:

$r_i^{(k)} + \sum_j P_{ij}^{(k)} v_j(1)$, to maximize $v(2)$ we will need $v^*(1)$

$$v_i^*(2) = \max_k \{ r_i^{(k)} + \sum_j P_{ij}^{(k)} v_j^*(1) \}$$

With n periods to go:

$$v_i^*(n) = \max_k \{ r_i^{(k)} + \sum_j P_{ij}^{(k)} v_j^*(n-1) \}$$

In vector format we have:

$$\mathbf{v}^*(n) = \max_A \{\mathbf{r}^A + [P^A]\mathbf{v}^*(n-1)\}$$

or equivalently:

$$\mathbf{v}^*(n) = \mathbf{r}^B + [P^B]\mathbf{v}^*(n-1)$$

With B such that

$$\mathbf{r}^B + [P^B]\mathbf{v}^*(n-1) \geq \mathbf{r}^A + [P^A]\mathbf{v}^*(n-1) \quad \forall A$$

Where A and B are policies (i.e. a decision k_i for each state i)

Optimal stationary policy

We assume for any policy A the resulting Markov chain with matrix $[P^A]$ is an ergodic unichain.

We know that for a policy A

$$\begin{aligned}\mathbf{w}^A &= \mathbf{r}^A - g^A \mathbf{e} + [P^A] \mathbf{w}^A \\&= \mathbf{r}^A - g^A \mathbf{e} + [P^A](\mathbf{r}^A - g^A \mathbf{e} + [P^A] \mathbf{w}^A) \\&= \mathbf{r}^A + [P^A] \mathbf{r}^A - 2g^A \mathbf{e} + [P^A]^2 \mathbf{w}^A \\&= \mathbf{r}^A + [P^A] \mathbf{r}^A + [P^A]^2 \mathbf{r}^A - 3g^A \mathbf{e} + [P^A]^3 \mathbf{w}^A \\&= \mathbf{r}^A + [P^A] \mathbf{r}^A + [P^A]^2 \mathbf{r}^A + \cdots + [P^A]^{n-1} \mathbf{r}^A - ng^A \mathbf{e} + [P^A]^n \mathbf{w}^A \\&\mathbf{v}^A(n) = \mathbf{r}^A + [P^A] \mathbf{r}^A + [P^A]^2 \mathbf{r}^A + \cdots + [P^A]^{n-1} \mathbf{r}^A + [P^A]^n \mathbf{v}(0)\end{aligned}$$

Hence:

$$\mathbf{v}^A(n) = ng^A \mathbf{e} + \mathbf{w}^A + [P^A]^n (\mathbf{v}(0) - \mathbf{w}^A)$$

The policy A that maximizes g^A will be called the optimal stationary policy.

A particular case

If $\mathbf{v}(0) = \mathbf{w}^B$ for a policy B such that

$$\mathbf{r}^B + [P^B]\mathbf{w}^B \geq \mathbf{r}^A + [P^A]\mathbf{w}^B \quad \forall A$$

then B is the dynamic optimal policy for each time period and

$$\mathbf{v}^*(n) = \mathbf{w}^B + ng^B\mathbf{e}$$

Proof:

$$\mathbf{v}^*(1) = \mathbf{r}^B + [P^B]\mathbf{w}^B = \mathbf{w}^B + g^B\mathbf{e} \quad \Rightarrow \mathbf{w}^B = \mathbf{v}^*(1) - g^B\mathbf{e}$$

Our hypothesis implies thus that:

$$\mathbf{r}^B + [P^B]\mathbf{v}^*(1) - g^B[P^B]\mathbf{e} \geq \mathbf{r}^A + [P^A]\mathbf{v}^*(1) - g^B[P^A]\mathbf{e} \quad \forall A$$

$$\Rightarrow \mathbf{r}^B + [P^B]\mathbf{v}^*(1) \geq \mathbf{r}^A + [P^A]\mathbf{v}^*(1) \quad \forall A$$

We can then repeat the same argument for $\mathbf{v}^*(2), \mathbf{v}^*(3), \dots$

Theorem

If for any policy A the Markov chain $[P^A]$ is an ergodic unichain, then the policy B is an optimal stationary policy if and only if:

$$\mathbf{r}^B + [P^B]\mathbf{w}^B \geq \mathbf{r}^A + [P^A]\mathbf{w}^B \quad \forall A \quad (*)$$

Proof:

\Leftarrow

If w^B were the final reward, then B would be the optimal dynamic policy.

If there existed a policy A such that $g^A > g^B$ then in the long term the policy A would generate a larger reward, hence a contradiction.

\Rightarrow

Let us assume $(*)$ is not satisfied.

Let $B = (k_1, k_2, \dots, k_J)$, there must exist i and k such that

$$r_i^{(k_i)} + \sum_j P_{ij}^{(k_i)} w_j^B \leq r_i^{(k)} + \sum_j P_{ij}^{(k)} w_j^B$$

Let us build $A = (k'_1, k'_2, \dots, k'_J)$ such that k'_i maximizes

$$r_i^{(k)} + \sum_j P_{ij}^{(k)} w_j^B$$

By our assumption there exists at least one index i such that $k_i \neq k'_i$. As a result:

$$\mathbf{r}^B + [P^B] \mathbf{w}^B \stackrel{\neq}{\leq} \mathbf{r}^A + [P^A] \mathbf{w}^B$$

$$\mathbf{w}^B + g^B \mathbf{e} \stackrel{\neq}{\leq} \mathbf{r}^A + [P^A] \mathbf{w}^B$$

Hence:

$$\boldsymbol{\pi}^A \mathbf{w}^B + g^B \boldsymbol{\pi}^A \mathbf{e} < \boldsymbol{\pi}^A \mathbf{r}^A + \boldsymbol{\pi}^A [P^A] \mathbf{w}^B \quad \Rightarrow \quad g^B < g^A$$

□

Policy Improvement algorithm

1. Choose an arbitrary policy B
2. Compute w^B
3. If $\mathbf{r}^B + [P^B]\mathbf{w}^B \geq \mathbf{r}^A + [P^A]\mathbf{w}^B \quad \forall A$ then stop
else continue
4. Find A such that $\mathbf{r}^A + [P^A]\mathbf{w}^B \not\geq \mathbf{r}^B + [P^B]\mathbf{w}^B$
5. $B \leftarrow A$ go to step 2.

Theorem

Assuming for any policy A the Markov chain $[P^A]$ is an ergodic unichain, if B is an optimal stationary policy then

$$\lim_{n \rightarrow \infty} \mathbf{v}^*(n) - n\mathbf{g}^B \mathbf{e} = \mathbf{w}^B + (\beta - \boldsymbol{\pi}^B \mathbf{w}^B) \mathbf{e}$$

- ▶ Average reward from optimal dynamic policy = \mathbf{g}^B
- ▶ Relative gain vector for optimal dynamic policy = \mathbf{w}^B
- ▶ Relative gain from all optimal stationary policies are identical
- ▶ $\min_i [v_i^*(n) - v_i^*(n-1)] \leq g^B \leq \max_i [v_i^*(n) - v_i^*(n-1)]$

Dynamic Programming algorithm for the stationary optimal policy

1. Fix an arbitrary vector $v(0)$
2. Compute $v^*(n) = \max_A \{r^A + [P^A]v^*(n-1)\}$
3. Compute $l = \min_i [v_i^*(n) - v_i^*(n-1)]$ and
 $u = \max_i [v_i^*(n) - v_i^*(n-1)]$
4. If $l < u$ go to step 2.
5. A is the optimal stationary policy and the maximum stationary reward is $l = u = g^A$.

Inventory Management example

You are managing inventories for a clinical trial, for drugs A the data are the following

- ▶ Weekly demand can be 0, 1 or 2 with respective probabilities 0.3 0.5 and 0.2
- ▶ Ordering cost is 10 per order + 20 per kit
- ▶ Holding cost is 4 per kit and per week
- ▶ The management says you cannot store more than 3 kits
- ▶ If you run out of units and a patient has to be supplied with an emergency shipment, it costs 50 per kit

Given that you can order at the end of a week for the beginning of next week (but you do not know the demand of next week), what would be the optimal ordering policy?